



# Sorting Optimization

Adam Todd



# Table of Contents

Background

Experiment

Analysis

Conclusions



# Background

- Sorting is an essential process for many applications
- Data Analytics
- E-Commerce
- Database querying



# Background

- Not all sorting algorithms lend nicely to multithreading
  - Bubble Sort
  - Insertion Sort
- To be efficient, we must use algorithms that are parallelizable
  - Merge sort
  - Quick Sort
- For top level performance, we must go beyond this into algorithms that are extremely parallelizable, and utilize GPU parallelism
  - Radix Sort
  - Bitonic Sort



# Experiment

To determine the effectiveness of multithreading on sorting algorithms, I implemented multiple sorting algorithms with different levels of parallelization.

- Single Threaded Merge Sort
- Multithreaded Merge Sort
- Radix Sort utilizing CUDA

Size	10	100	1000	10000	100000	1000000	10000000	100000000	1000000000
merge	0.000015	0.000131	0.001607	0.023406	<b>0.285668</b>	3.694614	47.471729	<b>595.80066</b>	
mt_merge	0.127111	0.135648	0.13781	0.226952	<b>0.36465</b>	2.345182	27.087753	<b>333.03562</b>	
radix	0.154003	0.14	0.141004	0.151	0.161	0.155999	0.240002	0.936044	60.263844
native	0.000004	0.000011	0.00011	0.001164	0.01812	0.363503	<b>4.99329</b>	57.720335	<b>837.217</b>
cupy native	<b>0.001</b>	<b>0.001</b>	<b>0.001</b>	<b>0.001</b>	<b>0.001</b>	<b>0.001</b>	<b>0.006</b>	<b>0.042</b>	<b>2.2886</b>

Experimental Data



# Analysis

- As expected, GPU level parallelism performed extremely well
- The already optimized python sort implementation was 38000% slower than the cupy sort for 1 billion item lists
- Even my own implementation of radix sort significantly outperformed the python implementation
- Significant overhead for small sizes, 4 million % slowdown for sorting 10 item list



# Analysis

- Out of memory issues...
  - Less optimal in place algorithm might be optimal depending on resource availability or speed requirements
- Power Consumption
  - GPU focused algorithms take a lot more power





# Conclusions

- For smaller list sizes, overhead typically worsens performance
- GPU parallelism can greatly improve performance
- Best algorithm heavily depends on specific use case, more than just best Big O
- When possible, keep data pre-sorted